

Forward School

Program Code: J620-002-4:2020

Program Name: FRONT-END SOFTWARE DEVELOPMENT

Title : Exe31 - MNIST Handwriting Exercise

Name: Ooi Caaron

IC Number: 990701-07-5837

Date :2/8/23

Introduction : Learning MNIST digit classification

Conclusion :

The Problem: MNIST digit classification

We're going to tackle a classic machine learning problem: MNIST handwritten digit classification. It's simple: given an image, classify it as a digit.

![[image.png]](attachment:image.png)

Each image in the MNIST dataset is 28x28 and contains a centered, grayscale digit. We'll flatten each 28x28 into a 784 dimensional vector, which we'll use as input to our neural network. Our output will be one of 10 possible classes: one for each digit.

Please check that you have the following packages installed (via conda or pip)
keras tensorflow numpy mnist

1. Setup

In [2]:

```
!pip install mnist
```

Collecting mnist

Downloading mnist-0.2.2-py2.py3-none-any.whl (3.5 kB)

Requirement already satisfied: numpy in c:\users\user\anaconda3\envs\python-ds-course\lib\site-packages (from mnist) (1.24.3)

Installing collected packages: mnist

Successfully installed mnist-0.2.2

In [3]:

```
#import all the required libraries
```

```
import numpy as np
```

```
import mnist
```

```
import keras
```

```
# The first time you run this might be a bit slow, since the
```

```
# mnist package has to download and cache the data.
```

```
train_images = mnist.train_images()
```

```
train_labels = mnist.train_labels()
```

```
test_images = mnist.test_images()
```

```
test_labels = mnist.test_labels()
```

Q: What's the dimension of the images data?

In [4]:

```
train_images.shape
```

Out[4]:

```
(60000, 28, 28)
```

In [6]:

```
test_images.shape
```

Out[6]:

```
(10000, 28, 28)
```

Q: What's the dimension of the label data?

In [7]:

```
train_labels.size
```

Out[7]:

```
60000
```

In [8]:

```
test_labels.size
```

Out[8]:

10000

Note: Curious about the dataset? try the following code. You can play around with the `image_index` value.

In [10]:

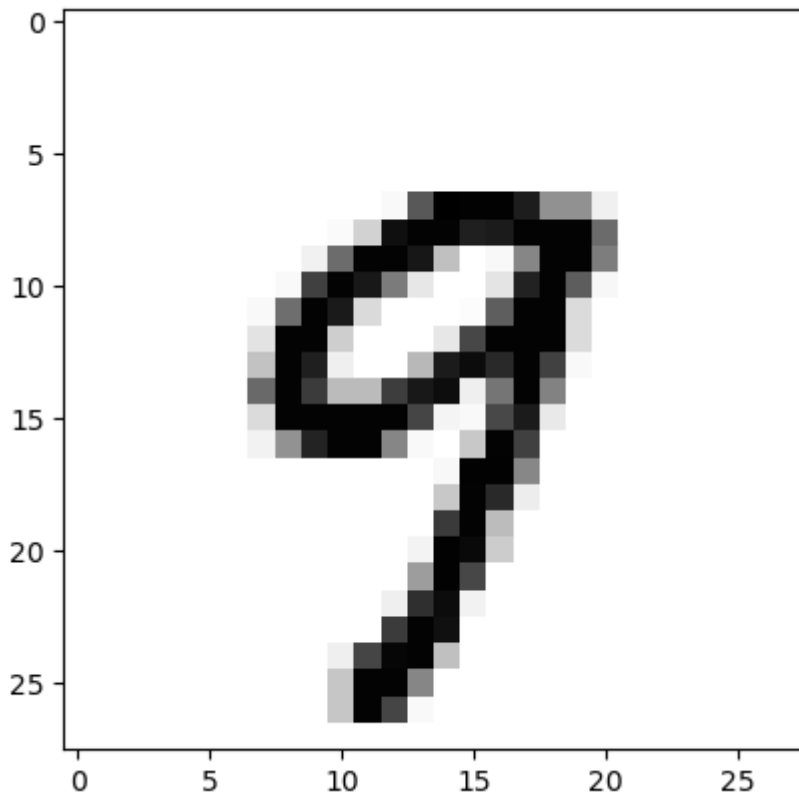
```
import matplotlib.pyplot as plt

image_index = 45 # You may select anything up to 60,000

print(train_labels[image_index]) # The label is 8
print(train_images[image_index])

plt.imshow(train_images[image_index], cmap='Greys')
```

```
<matplotlib.image.AxesImage at 0x1bf63a2d4d0>
```



2. Preparing the Data

As mentioned earlier, we need to flatten each image before we can pass it into our neural network. We'll also normalize the pixel values from $[0, 255]$ to $[-0.5, 0.5]$ to make our network easier to train (using smaller, centered values is often better).

In [16]:

```
# Normalize the images.
train_images = (train_images / 255) - 0.5
test_images = (test_images / 255) - 0.5

# Flatten the images.
train_images = train_images.reshape((-1, 784))
test_images = test_images.reshape((-1, 784))
```

Q: What's the dimension of the training and test images data?

In [11]:

```
train_images.shape
```

Out[11]:

```
(60000, 28, 28)
```

In [12]:

```
test_images.shape
```

Out[12]:

```
(10000, 28, 28)
```

3. Building the Model

Every Keras model is either built using the Sequential class, which represents a linear stack of layers, or the functional Model class, which is more customizable. We'll be using the simpler Sequential model, since our network is indeed a linear stack of layers.

Step: Start by instantiating a Sequential model.

- The first two layers have 64 nodes each and use the ReLU activation function.
- The last layer is a Softmax output layer with 10 nodes, one for each class.

Q: what's the correct input shape for your input layer?

In [17]:

```
from keras.models import Sequential
from keras.layers import Dense

# Define the model
model = Sequential([
    Dense(64, activation='relu'),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax'),
])
```

4. Compiling the Model

Before we can begin training, we need to configure the training process. We decide 3 key factors during the compilation step:

- The optimizer. We'll stick with a pretty good default: the Adam gradient-based optimizer. Keras has many other optimizers you can look into as well.
- The loss function. Since we're using a Softmax output layer, we'll use the Cross-Entropy loss. Keras distinguishes between `binary_crossentropy` (2 classes) and `categorical_crossentropy` (>2 classes), so we'll use the latter
- A list of metrics. Since this is a classification problem, we'll just have Keras report on the accuracy metric.

Step: Compile the model using the above options - adam, categorical_crossentropy, accuracy as metrics

In [18]:

```
model.compile(  
    optimizer='adam',  
    loss='categorical_crossentropy',  
    metrics=['accuracy'],  
)
```

5. Training the Model

Training a model in Keras literally consists only of calling `fit()` and specifying some parameters. There are a lot of possible parameters, but we'll only manually supply a few:

- The training data (images and labels), commonly known as X and Y, respectively.
- The number of epochs (iterations over the entire dataset) to train for.
- The batch size (number of samples per gradient update) to use when training.

Step: set epochs to a suitable number, and `batch_size = 32`

In [19]:

```
from keras.models import Sequential  
from keras.layers import Dense  
from keras.utils import to_categorical  
  
# Train the model.  
model.fit(  
    train_images,  
    to_categorical(train_labels),  
    epochs=5,  
    batch_size=32,  
)
```

```
Epoch 1/5  
1875/1875 [=====] - 4s 2ms/step - loss: 0.3797 -  
accuracy: 0.8857  
Epoch 2/5  
1875/1875 [=====] - 3s 2ms/step - loss: 0.2054 -  
accuracy: 0.9366  
Epoch 3/5  
1875/1875 [=====] - 3s 2ms/step - loss: 0.1571 -  
accuracy: 0.9527  
Epoch 4/5  
1875/1875 [=====] - 3s 2ms/step - loss: 0.1322 -  
accuracy: 0.9593  
Epoch 5/5  
1875/1875 [=====] - 3s 2ms/step - loss: 0.1174 -  
accuracy: 0.9633
```

Out[19]:

```
<keras.src.callbacks.History at 0x1bf63b64090>
```

Q: Do you run into any problem? Why?

In []:

Q: what's your achieved accuracy?

6. Testing the Model

Step: Evaluating the model by testing against the test data

In [20]:

```
model.evaluate(
    test_images,
    to_categorical(test_labels)
)
```

313/313 [=====] - 1s 2ms/step - loss: 0.1403 - ac
curacy: 0.9573

Out[20]:

```
[0.14030496776103973, 0.9573000073432922]
```

7. Using the Model

Now that we have a working, trained model, let's put it to use. The first thing we'll do is save it to disk so we can load it back up anytime.

Step: save the model using the save_weights function

In [16]:

```
model.save_weights('model.h5')
```

In [17]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Build the model.
model = Sequential([
    Dense(64, activation='relu', input_shape=(784,)),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax'),
])

# Load the model's saved weights.
model.load_weights('model.h5')
```

8. Predict

Using the trained model to make predictions is easy: we pass an array of inputs to `predict()` and it returns an array of outputs. Keep in mind that the output of our network is 10 probabilities (because of softmax), so we'll use `np.argmax()` to turn those into actual digits.

In [18]:

```
# Predict on the first 5 test images.
predictions = model.predict(test_images[:5])

# Print our model's predictions.
print(np.argmax(predictions, axis=1)) # [7, 2, 1, 0, 4]

# Check our predictions against the ground truths.
print(test_labels[:5]) # [7, 2, 1, 0, 4]
```

```
[7 2 1 0 4]
[7 2 1 0 4]
```

Note: What's the difference between `model.save_weights` and `model.save`? -

[https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save\(\)%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else](https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save()%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else)
[.https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save\(\)%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else\).](https://stackoverflow.com/questions/42621864/difference-between-keras-model-save-and-model-save-weights#:~:text=save()%20saves%20the%20weights,to%20HDF5%20and%20nothing%20else)

This exercise is adapted from <https://victorzhou.com/blog/keras-neural-network-tutorial/>
[\(https://victorzhou.com/blog/keras-neural-network-tutorial/\)](https://victorzhou.com/blog/keras-neural-network-tutorial/)

Challenge 1:

Retrain your model by using different network depths - what will you conclude?

In []:

Challenge 2:

Retrain your model by using different activation (other than ReLU) - what differences does it make?

In [21]:

```
model = Sequential([
    Dense(64, activation='sigmoid', input_shape=(784,)),
    Dense(64, activation='sigmoid'),
    Dense(10, activation='softmax'),
])
```

Challenge 3:

Fit your model using `validation_data` option - What differences will that bring?

In [22]:

```
from tensorflow.keras.optimizers import Adam

model.compile(
    optimizer=Adam(lr=0.005),
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

model.fit(
    train_images,
    to_categorical(train_labels),
    epochs=5,
    batch_size=32,
    validation_data=(test_images, to_categorical(test_labels))
)
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.Adam`.

```
Epoch 1/5
1875/1875 [=====] - 6s 3ms/step - loss: 0.5660 -
accuracy: 0.8584 - val_loss: 0.2610 - val_accuracy: 0.9253
Epoch 2/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.2313 -
accuracy: 0.9316 - val_loss: 0.1967 - val_accuracy: 0.9418
Epoch 3/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1713 -
accuracy: 0.9491 - val_loss: 0.1673 - val_accuracy: 0.9494
Epoch 4/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1349 -
accuracy: 0.9599 - val_loss: 0.1424 - val_accuracy: 0.9576
Epoch 5/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1141 -
accuracy: 0.9652 - val_loss: 0.1301 - val_accuracy: 0.9597
```

Out[22]:

```
<keras.src.callbacks.History at 0x1bf62ebacd0>
```

Challenge 4:

How will you load your saved weights to use it in a separate code? Upload your saved model/weights, and compare your model/weights with a model/weights from one of your classmate's.

In []:

```
model.save_weights('model.h5')
```

Challenge 5:

How can you load any image from the data set and let your model (or your classmate's) to predict the image?

In [23]:

```
predictions = model.predict(test_images[:5])  
print(np.argmax(predictions, axis = 1))  
print(test_labels[:5])
```

```
1/1 [=====] - 0s 124ms/step  
[7 2 1 0 4]  
[7 2 1 0 4]
```

In []: